

# Parallel Numerical Algorithms

## Chapter 6 – LU Factorization

Prof. Michael T. Heath

Department of Computer Science  
University of Illinois at Urbana-Champaign

CS 554 / CSE 512

Michael T. Heath

Parallel Numerical Algorithms

1 / 42

LU Factorization  
Parallel Algorithms for LU  
Partial PivotingMotivation  
Gaussian Elimination

## LU Factorization

- System of linear algebraic equations has form

$$Ax = b$$

where  $A$  is given  $n \times n$  matrix,  $b$  is given  $n$ -vector, and  $x$  is unknown solution  $n$ -vector to be computed

- Direct method for solving general linear system is by computing **LU factorization**

$$A = LU$$

where  $L$  is unit lower triangular and  $U$  is upper triangular

Michael T. Heath

Parallel Numerical Algorithms

2 / 42

LU Factorization  
Parallel Algorithms for LU  
Partial PivotingMotivation  
Gaussian Elimination

## Gaussian Elimination Algorithm

LU factorization can be computed by Gaussian elimination as follows, where  $U$  overwrites  $A$

```

for  $k = 1$  to  $n - 1$            { loop over columns }
  for  $i = k + 1$  to  $n$          { compute multipliers
     $\ell_{ik} = a_{ik}/a_{kk}$          for current column }
  end
  for  $j = k + 1$  to  $n$ 
    for  $i = k + 1$  to  $n$        { apply transformation to
       $a_{ij} = a_{ij} - \ell_{ik}a_{kj}$  remaining submatrix }
    end
  end
end
end

```

Michael T. Heath

Parallel Numerical Algorithms

5 / 42

LU Factorization  
Parallel Algorithms for LU  
Partial PivotingMotivation  
Gaussian Elimination

## Loop Orderings for Gaussian Elimination

- Gaussian elimination has general form of triple-nested loop in which entries of  $L$  and  $U$  overwrite those of  $A$

```

for _____
  for _____
    for _____
       $a_{ij} = a_{ij} - (a_{ik}/a_{kk})a_{kj}$ 
    end
  end
end
end

```

- Indices  $i$ ,  $j$ , and  $k$  of **for** loops can be taken in any order, for total of  $3! = 6$  different ways of arranging loops

Michael T. Heath

Parallel Numerical Algorithms

7 / 42

## Outline

- 1 LU Factorization
  - Motivation
  - Gaussian Elimination
- 2 Parallel Algorithms for LU
  - Fine-Grain Algorithm
  - Agglomeration Schemes
  - Scalability
- 3 Partial Pivoting

Michael T. Heath

Parallel Numerical Algorithms

2 / 42

LU Factorization  
Parallel Algorithms for LU  
Partial PivotingMotivation  
Gaussian Elimination

## LU Factorization

- System  $Ax = b$  then becomes

$$LUx = b$$

- Solve lower triangular system

$$Ly = b$$

by forward-substitution to obtain vector  $y$

- Finally, solve upper triangular system

$$Ux = y$$

by back-substitution to obtain solution  $x$  to original system

Michael T. Heath

Parallel Numerical Algorithms

4 / 42

LU Factorization  
Parallel Algorithms for LU  
Partial PivotingMotivation  
Gaussian Elimination

## Gaussian Elimination Algorithm

- In general, row interchanges (pivoting) may be required to ensure existence of LU factorization and numerical stability of Gaussian elimination algorithm, but for simplicity we temporarily ignore this issue
- Gaussian elimination requires about  $n^3/3$  paired additions and multiplications, so model serial time as

$$T_1 = t_c n^3/3$$

where  $t_c$  is time required for multiply-add operation

- About  $n^2/2$  divisions also required, but we ignore this lower-order term

Michael T. Heath

Parallel Numerical Algorithms

6 / 42

LU Factorization  
Parallel Algorithms for LU  
Partial PivotingMotivation  
Gaussian Elimination

## Loop Orderings for Gaussian Elimination

- Different loop orders have different memory access patterns, which may cause their performance to vary widely, depending on architectural features such as cache, paging, vector registers, etc.
- Perhaps most promising for parallel implementation are  $kij$  and  $kji$  forms, which differ only in accessing matrix by rows or columns, respectively

Michael T. Heath

Parallel Numerical Algorithms

8 / 42

## Gaussian Elimination Algorithm

- $kji$  form of Gaussian elimination

```


for  $k = 1$  to  $n - 1$ 
  for  $i = k + 1$  to  $n$ 
     $\ell_{ik} = a_{ik} / a_{kk}$ 
  end
  for  $j = k + 1$  to  $n$ 
    for  $i = k + 1$  to  $n$ 
       $a_{ij} = a_{ij} - \ell_{ik} a_{kj}$ 
    end
  end
end
end

```

- Multipliers  $\ell_{ik}$  computed outside inner loop for greater efficiency

Michael T. Heath Parallel Numerical Algorithms 9 / 42

## Fine-Grain Tasks and Communication



Michael T. Heath Parallel Numerical Algorithms 10 / 42

## Agglomeration


### Agglomerate

With  $n \times n$  array of fine-grain tasks, natural strategies are

- 2-D: combine  $k \times k$  subarray of fine-grain tasks to form each coarse-grain task, yielding  $(n/k)^2$  coarse-grain tasks
- 1-D column: combine  $n$  fine-grain tasks in each column into coarse-grain task, yielding  $n$  coarse-grain tasks
- 1-D row: combine  $n$  fine-grain tasks in each row into coarse-grain task, yielding  $n$  coarse-grain tasks

Michael T. Heath Parallel Numerical Algorithms 13 / 42

## 1-D Column Agglomeration



Michael T. Heath Parallel Numerical Algorithms 15 / 42

## Parallel Algorithm

### Partition

- For  $i, j = 1, \dots, n$ , fine-grain task  $(i, j)$  stores  $a_{ij}$  and computes and stores

$$\begin{cases} u_{ij}, & \text{if } i \leq j \\ \ell_{ij}, & \text{if } i > j \end{cases}$$

yielding 2-D array of  $n^2$  fine-grain tasks

### Communicate

- Broadcast entries of  $A$  vertically to tasks below
- Broadcast entries of  $L$  horizontally to tasks to right

Michael T. Heath Parallel Numerical Algorithms 10 / 42

## Fine-Grain Parallel Algorithm


```

for  $k = 1$  to  $\min(i, j) - 1$ 
  rcv broadcast of  $a_{kj}$  from task  $(k, j)$  { vert bcast }
  rcv broadcast of  $\ell_{ik}$  from task  $(i, k)$  { horiz bcast }
   $a_{ij} = a_{ij} - \ell_{ik} a_{kj}$  { update entry }
end
if  $i \leq j$  then
  broadcast  $a_{ij}$  to tasks  $(k, j)$ ,  $k = i + 1, \dots, n$  { vert bcast }
else
  rcv broadcast of  $a_{jj}$  from task  $(i, j)$  { vert bcast }
   $\ell_{ij} = a_{ij} / a_{jj}$  { multiplier }
  broadcast  $\ell_{ij}$  to tasks  $(i, k)$ ,  $k = j + 1, \dots, n$  { horiz bcast }
end

```


Michael T. Heath Parallel Numerical Algorithms 12 / 42

## 2-D Agglomeration



Michael T. Heath Parallel Numerical Algorithms 14 / 42

## 1-D Row Agglomeration



Michael T. Heath Parallel Numerical Algorithms 16 / 42

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

## Mapping

## Map

- 2-D: assign  $(n/k)^2/p$  coarse-grain tasks to each of  $p$  processes using any desired mapping in each dimension, treating target network as 2-D mesh
- 1-D: assign  $n/p$  coarse-grain tasks to each of  $p$  processes using any desired mapping, treating target network as 1-D mesh

Michael T. Heath

Parallel Numerical Algorithms

17 / 42

## Coarse-Grain 2-D Parallel Algorithm

```

for  $k = 1$  to  $n - 1$ 
  broadcast  $\{a_{kj} : j \in \text{mycols}, j \geq k\}$  in process column
  if  $k \in \text{mycols}$  then
    for  $i \in \text{myrows}, i > k$ 
       $\ell_{ik} = a_{ik}/a_{kk}$  { multipliers }
    end
  end
  broadcast  $\{\ell_{ik} : i \in \text{myrows}, i > k\}$  in process row
  for  $j \in \text{mycols}, j > k$ 
    for  $i \in \text{myrows}, i > k$ 
       $a_{ij} = a_{ij} - \ell_{ik} a_{kj}$  { update }
    end
  end
end
end

```

Michael T. Heath

Parallel Numerical Algorithms

18 / 42

## Performance Enhancements

Performance can also be enhanced by overlapping communication and computation

- At step  $k$ , each process completes updating its portion of remaining unreduced submatrix before moving on to step  $k + 1$
- Broadcast of each segment of row  $k + 1$ , and computation and broadcast of each segment of multipliers for step  $k + 1$ , could be initiated as soon as relevant segments of row  $k + 1$  and column  $k + 1$  have been updated by their owners, before completing remainder of their updating for step  $k$
- This *send ahead* strategy enables other processes to start working on next step earlier than they otherwise could

Michael T. Heath

Parallel Numerical Algorithms

21 / 42

## 1-D Column Agglomeration


- Matrix rows need not be broadcast vertically, since any given column is contained entirely in only one process
- But there is no parallelism in computing multipliers or updating any given column
- Horizontal broadcasts still required to communicate multipliers for updating

Michael T. Heath

Parallel Numerical Algorithms

23 / 42

## 2-D Agglomeration with Cyclic Mapping



Michael T. Heath

Parallel Numerical Algorithms

18 / 42

## Performance Enhancements


- Each process becomes idle as soon as its last row and column are completed
- With block mapping, in which each process holds contiguous block of rows and columns, some processes become idle long before overall computation is complete
- Block mapping also yields unbalanced load, as computing multipliers and updates requires successively less work with increasing row and column numbers
- Cyclic or reflection mapping improves both concurrency and load balance

Michael T. Heath

Parallel Numerical Algorithms

20 / 42

## 1-D Column Agglomeration with Cyclic Mapping



Michael T. Heath

Parallel Numerical Algorithms

22 / 42

## Coarse-Grain 1-D Column Parallel Algorithm

```

for  $k = 1$  to  $n - 1$ 
  if  $k \in \text{mycols}$  then
    for  $i = k + 1$  to  $n$ 
       $\ell_{ik} = a_{ik}/a_{kk}$  { multipliers }
    end
  end
  broadcast  $\{\ell_{ik} : k < i \leq n\}$  { broadcast }
  for  $j \in \text{mycols}, j > k$ 
    for  $i = k + 1$  to  $n$ 
       $a_{ij} = a_{ij} - \ell_{ik} a_{kj}$  { update }
    end
  end
end
end


```

Michael T. Heath

Parallel Numerical Algorithms

24 / 42

## 1-D Row Agglomeration with Cyclic Mapping



Michael T. Heath Parallel Numerical Algorithms 25 / 42

## Coarse-Grain 1-D Row Parallel Algorithm

```

for k = 1 to n - 1
  broadcast {akj : k ≤ j ≤ n}    { broadcast }
  for i ∈ myrows, i > k,
    ℓik = aik/akk                { multipliers }
  end
  for j = k + 1 to n
    for i ∈ myrows, i > k,
      aij = aij - ℓikakj        { update }
    end
  end
end
end

```

Michael T. Heath Parallel Numerical Algorithms 25 / 42

## Scalability for 2-D Agglomeration

- Updating by each process at step  $k$  requires about  $(n - k)^2/p$  operations
- Summing over  $n - 1$  steps

$$\begin{aligned}
 T_{\text{comp}} &\approx t_c \sum_{k=1}^{n-1} (n - k)^2/p \\
 &\approx t_c n^3/(3p)
 \end{aligned}$$

Michael T. Heath Parallel Numerical Algorithms 29 / 42

## Scalability for 2-D Agglomeration

- Total execution time is

$$T_p \approx t_c n^3/(3p) + 2t_s n + t_w n^2/\sqrt{p}$$

- To determine isoefficiency function, set

$$t_c n^3/3 \approx E(t_c n^3/3 + 2t_s n p + t_w n^2 \sqrt{p})$$

which holds for large  $p$  if  $n = \Theta(\sqrt{p})$ , so isoefficiency function is  $\Theta(p\sqrt{p})$ , since  $T_1 = \Theta(n^3)$

Michael T. Heath Parallel Numerical Algorithms 31 / 42

## 1-D Row Agglomeration

- Multipliers need not be broadcast horizontally, since any given matrix row is contained entirely in only one process
- But there is no parallelism in updating any given row
- Vertical broadcasts still required to communicate each row of matrix to processes below it for updating

Michael T. Heath Parallel Numerical Algorithms 26 / 42

## Performance Enhancements

- Same performance enhancements as for 2-D agglomeration apply to both 1-D column and 1-D row agglomerations as well, including cyclic mapping and send ahead strategy

Michael T. Heath Parallel Numerical Algorithms 28 / 42

## Scalability for 2-D Agglomeration

- Similarly, amount of data broadcast at step  $k$  along each process row and column is about  $(n - k)/\sqrt{p}$ , so on 2-D mesh

$$\begin{aligned}
 T_{\text{comm}} &\approx \sum_{k=1}^{n-1} 2(t_s + t_w(n - k)/\sqrt{p}) \\
 &\approx 2t_s n + t_w n^2/\sqrt{p}
 \end{aligned}$$

where we have allowed for overlap of broadcasts for successive steps

Michael T. Heath Parallel Numerical Algorithms 30 / 42

## Scalability for 1-D Agglomeration

- With either 1-D column or 1-D row agglomeration, updating by each process at step  $k$  requires about  $(n - k)^2/p$  operations
- Summing over  $n - 1$  steps

$$\begin{aligned}
 T_{\text{comp}} &\approx t_c \sum_{k=1}^{n-1} (n - k)^2/p \\
 &\approx t_c n^3/(3p)
 \end{aligned}$$

Michael T. Heath Parallel Numerical Algorithms 32 / 42

## Scalability for 1-D Agglomeration

- Amount of data broadcast at step  $k$  is about  $n - k$ , so on 1-D mesh

$$T_{\text{comm}} \approx \sum_{k=1}^{n-1} (t_s + t_w (n - k))$$

$$\approx t_s n + t_w n^2 / 2$$

where we have allowed for overlap of broadcasts for successive steps

Michael T. Heath

Parallel Numerical Algorithms

33 / 42

## Partial Pivoting

- Row ordering of  $A$  is irrelevant in system of linear equations
- Partial pivoting takes rows in order of largest entry in magnitude of leading column of remaining unreduced matrix
- This choice ensures that multipliers do not exceed 1 in magnitude, which reduces amplification of rounding errors
- In general, partial pivoting is required to ensure existence and numerical stability of LU factorization

Michael T. Heath

Parallel Numerical Algorithms

34 / 42

## Parallel Partial Pivoting

- Partial pivoting complicates parallel implementation of Gaussian elimination and significantly affects potential performance
- With 2-D algorithm, pivot search is parallel but requires communication within process column and inhibits overlapping of successive steps
- With 1-D column algorithm, pivot search requires no communication but is purely serial
- Once pivot is found, index of pivot row must be communicated to other processes, and rows must be explicitly or implicitly interchanged in each process

Michael T. Heath

Parallel Numerical Algorithms

37 / 42

## Alternatives to Partial Pivoting

- Because of negative effects of partial pivoting on parallel performance, various alternatives have been proposed that limit pivot search
  - tournament pivoting
  - threshold pivoting
  - pairwise pivoting
- Such strategies are not foolproof and may trade off some degree of stability and accuracy for speed
- Stability and accuracy may be recovered via iterative refinement, but this has its own cost

Michael T. Heath

Parallel Numerical Algorithms

39 / 42

## Scalability for 1-D Agglomeration

- Total execution time is

$$T_p \approx t_c n^3 / (3p) + t_s n + t_w n^2 / 2$$

- To determine isoefficiency function, set

$$t_c n^3 / 3 \approx E (t_c n^3 / 3 + t_s n p + t_w n^2 p / 2)$$

which holds for large  $p$  if  $n = \Theta(p)$ , so isoefficiency function is  $\Theta(p^3)$ , since  $T_1 = \Theta(n^3)$

Michael T. Heath

Parallel Numerical Algorithms

34 / 42

## Partial Pivoting

- Partial pivoting yields factorization of form

$$PA = LU$$

where  $P$  is permutation matrix

- If  $PA = LU$ , then system  $Ax = b$  becomes

$$PAx = LUx = Pb$$

which can be solved by forward-substitution in lower triangular system  $Ly = Pb$ , followed by back-substitution in upper triangular system  $Ux = y$

Michael T. Heath

Parallel Numerical Algorithms

36 / 42

## Parallel Partial Pivoting

- With 1-D row algorithm, pivot search is parallel but requires communication among processes and inhibits overlapping of successive steps
- If rows are explicitly interchanged, then only two processes are involved
- If rows are implicitly interchanged, then mapping of rows to processes is altered, which may degrade concurrency and load balance
- Tradeoff between column and row algorithms with partial pivoting depends on relative speeds of communication and computation

Michael T. Heath

Parallel Numerical Algorithms

38 / 42

## Communication vs. Memory Tradeoff

- If explicit replication of storage is allowed, then lower communication volume is possible
- As with matrix multiplication, "2.5-D" algorithms have recently been developed that use partial storage replication to reduce communication volume to whatever extent available memory allows
- If sufficient memory is available, then these algorithms can achieve provably optimal communication

Michael T. Heath

Parallel Numerical Algorithms

40 / 42

## References

- J. W. Demmel, M. T. Heath, and H. A. van der Vorst, Parallel numerical linear algebra, *Acta Numerica* 2:111-197, 1993
- G. A. Geist and C. H. Romine, LU factorization algorithms on distributed-memory multiprocessor architectures, *SIAM J. Sci. Stat. Comput.* 9:639-649, 1988
- L. Grigori, J. Demmel, and H. Xiang, CALU: A communication optimal LU factorization algorithm, *SIAM J. Matrix Anal. Appl.* 32:1317-1350, 2011
- B. A. Hendrickson and D. E. Womble, The torus-wrap mapping for dense matrix calculations on massively parallel computers, *SIAM J. Sci. Stat. Comput.* 15:1201-1226, 1994

## References

- J. M. Ortega, *Introduction to Parallel and Vector Solution of Linear Systems*, Plenum Press, 1988
- J. M. Ortega and C. H. Romine, The *ijk* forms of factorization methods II: parallel systems, *Parallel Comput.* 7:149-162, 1988
- Y. Robert, *The Impact of Vector and Parallel Architectures on the Gaussian Elimination Algorithm*, Wiley, 1990
- E. Solomonik and J. Demmel, Communication-optimal parallel 2.5D matrix multiplication and LU factorization algorithms, 17th Euro-Par Conf. on Parallel Processing, LNCS 6853, Springer, 2011
- S. A. Vavasis, Gaussian elimination with pivoting is P-complete, *SIAM J. Disc. Math.* 2:413-423, 1989

# Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder